

APPROVAL SHEET

Title of Thesis: A Framework for Predicting and Controlling System-Level Properties of Agent-Based Models

Name of Candidate: Donald P. Miner
PhD in Computer Science,
2010

Thesis and Abstract Approved: _____
Dr. Marie desJardins
Associate Professor
Department of Computer Science and
Electrical Engineering

Date Approved: _____

Curriculum Vitae

Name: MY-FULL-NAME.

Permanent Address: MY-FULL-ADDRESS.

Degree and date to be conferred: DEGREE-NAME, GRADUATION-MONTH
GRADUATION-YEAR.

Date of Birth: MY-BIRTHDATE.

Place of Birth: MY-PLACE-OF-BIRTH.

Secondary Education: MY-HIGH-SCHOOL, MY-HIGH-SCHOOLS-CITY,
MY-HIGH-SCHOOLS-STATE.

Collegiate institutions attended:

University of Maryland Baltimore County, DEGREE-NAME MY-MAJOR,
GRADUATION-YEAR.
MY-OTHER-DEGREES.

Major: MY-MAJOR.

Minor: MY-MINOR.

Professional publications:

FULL-CITATION-INFORMATION.
FULL-CITATION-INFORMATION.

Professional positions held:

EMPLOYMENT-INFO. (START-DATE – END-DATE).
EMPLOYMENT-INFO. (START-DATE – END-DATE).

ABSTRACT

Title of Thesis: A Framework for Predicting and Controlling System-Level Properties of Agent-Based Models

Donald P. Miner, PhD in Computer Science, 2010

Thesis directed by: Dr. Marie desJardins, Associate Professor
Department of Computer Science and
Electrical Engineering

This is the abstract. Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum

**A Framework for Predicting and Controlling
System-Level Properties of Agent-Based Models**

by
Donald P. Miner

Thesis submitted to the Faculty of the Graduate School
of the University of Maryland in partial fulfillment
of the requirements for the degree of
Doctor of Philosophy in Computer Science
2010

This is my dedication.

ACKNOWLEDGMENTS

These will be written out later:

- John Way: My excellent high school computer science teacher; he was the first teacher to truly inspire and make me interested in something
- Richard Chang: Unofficial undergrad advisor; Inspired me to work on hard problems with his classes/my undergrad thesis; Indirectly helped me want to pursue graduate school
- Marie: advisor; introducing me to MAS
- Bill Rand: introducing me to NetLogo
- Forrest Stonedahl: working on a similar problem; a conversation which was a turning point in my research focus to ABMs; introduced me to the term 'meta-model'
- Tim Oates: Suggestions dealing with the ML portion of my research
- Undergraduate Researchers (Peter, Kevin, Doug, Nathan?): Helping with researching new domains
- Marc Pickett: Good friend that is always willing to listen to research ideas; Helped me throughout grad school
- Senior grad students who helped me as a young grad student: Adam Anthony, Eric Eaton, Blaz Bulka
- Other supporting CS graduate students: Wes Griffin, Yasaman Haghpanah, Niels Kasch, James MacGlashan, JC Montminy, Sourav M, Patti Ordonez, Soumi Ray, and Brandon Wilson.

TABLE OF CONTENTS

| | | |
|------------------------|---|-----------|
| DEDICATION | ii | |
| ACKNOWLEDGMENTS | iii | |
| LIST OF FIGURES | vii | |
| LIST OF TABLES | viii | |
| Chapter 1 | INTRODUCTION AND MOTIVATION | 1 |
| 1.1 | Agent-Based Models | 1 |
| 1.2 | Wolves, Sheep and Grass | 4 |
| 1.3 | Overview of the ABM Meta-Modeling Framework | 7 |
| 1.4 | Summary of Contributions | 8 |
| 1.5 | Dissertation Organization | 9 |
| Chapter 2 | THE ABM META-MODELING FRAMEWORK | 10 |
| 2.1 | Design Goals of The ABM Meta-Modeling Framework | 11 |
| 2.2 | Framework Structure | 12 |
| 2.2.1 | Defining System-Level Behavior Properties | 13 |
| 2.2.2 | Sampling | 15 |

| | | |
|------------------|--|-----------|
| 2.2.3 | The Forward-Mapping Problem | 16 |
| 2.2.4 | The Reverse-Mapping Problem | 17 |
| 2.2.5 | Summary of Configuration Points | 18 |
| 2.3 | Software Implementation Details | 19 |
| 2.4 | Analysis of AMF vs. the Design Goals | 20 |
| Chapter 3 | RELATED WORK | 22 |
| 3.1 | Experimentation in ABMs | 22 |
| 3.2 | Prediction of System-Level Behavior | 22 |
| 3.3 | Inversion of Neural Networks | 23 |
| Chapter 4 | BACKGROUND | 24 |
| 4.1 | Agent-Based Modeling | 24 |
| 4.1.1 | Examples of Agent-Based Models | 27 |
| 4.1.2 | Models of Human Societies | 29 |
| 4.1.3 | Multi-Agent Software Frameworks | 31 |
| 4.2 | Regression | 36 |
| 4.2.1 | K-Nearest Neighbor Regression | 38 |
| 4.2.2 | Robust Locally Weighted Regression and Smoothing Scatterplots (LOESS) | 39 |
| 4.2.3 | Nonlinear Regression | 40 |
| 4.2.4 | Multilinear Interpolation | 40 |
| 4.3 | Inverting Regression | 41 |
| 4.3.1 | Plane Intersection | 41 |
| 4.3.2 | Optimization | 41 |

| | | |
|-------------------|--|-----------|
| Chapter 5 | DEFINING SYSTEM-LEVEL PROPERTIES | 42 |
| Chapter 6 | THE FORWARD MAPPING PROBLEM | 44 |
| Chapter 7 | THE REVERSE MAPPING PROBLEM | 45 |
| Chapter 8 | USING META-MODELS | 46 |
| Chapter 9 | RESULTS | 47 |
| Chapter 10 | CONCLUSIONS AND FUTURE WORK | 48 |
| Appendix A | CODE SAMPLES | 49 |
| | A.1 Sample Java/NetLogo Sampling Program | 49 |
| Appendix B | INTERFACE SPECIFICATION FOR REGRESSION ALGO- RITHMS | 50 |
| | REFERENCES | 51 |

LIST OF FIGURES

| | | |
|-----|---|----|
| 1.1 | A screenshot of NetLogo's graphical user interface while executing a flocking simulation. | 2 |
| 1.2 | A screen shot from NetLogo's Wolf Sheep Predation model. | 4 |
| 1.3 | The control and monitor interface for the Wolf Sheep Predation model. . . . | 5 |
| 1.4 | Differences in populations based on changes of the <i>sheep-gain-from-food</i> parameter. | 6 |
| 2.1 | An overview of the phases of AMF and how data flows between them. . . . | 14 |
| 4.1 | A boid flock moving through a two-dimensional space. | 27 |

LIST OF TABLES

| | | |
|-----|--|----|
| 2.1 | Sample Predictions of Behavior in the Wolf Sheep Predation Model | 17 |
|-----|--|----|

Chapter 1

INTRODUCTION AND MOTIVATION

The behavior of individual agents in an agent-based model (ABM) is typically well understood because the agent's program directly controls its local behaviors. What is typically not understood is how changing these programs' agent-level control parameters affect the observed system-level behaviors of the ABM. The aim of this dissertation is to provide researchers and users of ABMs insight into how these agent-level parameters affect system-level properties. In this dissertation, I discuss a learning framework named the ABM Meta-Modeling Framework (AMF) that I have developed that can be used to predict and control system-level behaviors of agent-based models. With this framework, users can interact with ABMs in terms of intuitive system-level concepts, instead of with agent-level controls that only indirectly affect system-level behaviors.

1.1 Agent-Based Models

Agent-based models are used by scientists to analyze system-level behaviors of complex systems by simulating the system bottom-up. At the bottom of these simulations are individual agents that locally interact with other agents and the environment. All the behavior in an ABM, from agent-level local interactions to system-level behaviors, emerge from these local interactions, which are governed by the individual *agent programs*. ABMs can

be used to understand how changes in individuals' *agent-level parameters* affect *system-level properties*.




FIG. 1.1. A screenshot of NetLogo's graphical user interface while executing a flocking simulation.

Agent-level control parameters adjust the behaviors of agent-based programs. However, scientists are not typically interested in the local interactions between agents—they are interested in the resulting system-level behaviors that result. For example, researchers that studied agent-based models of lane formations in army ants were interested in the traffic patterns of the lanes, not the individual behaviors of the ants (Couzin & Franks 2003). In other work, researchers that studied locusts were interested in determining at what critical density locusts begin to swarm and destroy crops (Buhl *et al.* 2006). Typically, scientists analyze ABMs by viewing visualizations of the environment or gathering statistical data on the simulation. For instance, NetLogo, an agent-based modeling programming

environment (Tisue & Wilensky 2004), has monitors, plots and visualizations to convey system-level properties to the user. In Figure 1.1, monitors are displaying *density*, *average-heading* and *stddev-heading* statistics for a flocking domain. In addition, a plot of density shows how it has changed over time. These tools are used by a researcher to generate a mental model of how the agent-level control parameters of the flocking domain (the sliders seen in the user interface) affect these system-level properties.

Although using ABMs for researching agent-based systems has been proven useful in a number of domains, there is a glaring conceptual disconnect from the user's perspective, between the agent-level controls and the system-level properties. The classical ABM control method of adjusting agent-level properties is unintuitive because they only indirectly affect the system-level properties through emergence. With the current methodology, a simulation has to be executed in order to observe what the values of the system-level properties will be. A time consuming iterative process of guess-and-check is the only way to configure the system to have it exhibit a desired system-level behavior. A determination of what an ABM will do at a system-level, given only the agent-level parameters, is not possible with current software.

The main goal of the ABM Meta-Modeling Framework is to bridge the gap between agent-level parameters and system-level properties. AMF reduces the learning curve of an agent-based model since users are interacting with the system at the system-level, instead of at the agent-level. Qualitative analysis of an ABM's system-level properties will be a more efficient process since researchers deal with an abstraction of the system's controls. In addition, the models learned by AMF can be inspected to gather quantitative data about the correlations between system-level properties and agent-level parameters.




FIG. 1.2. A screen shot from NetLogo's Wolf Sheep Predation model.

1.2 Wolves, Sheep and Grass

Throughout this dissertation, I will use NetLogo's Wolf Sheep Predation model (Wilensky 1997c), which is bundled with NetLogo's standard Model Library,¹ as an example to explain concepts. A snapshot of its NetLogo visualization is shown in Figure 1.2. This multi-agent model simulates a food chain consisting of wolf agents, sheep agents and grass in a two-dimensional space. The model is controlled by seven agent-level control parameters, which directly affect the following agent behaviors:

- The system is initialized with *initial-number-sheep* sheep and *initial-number-wolves* wolves.
- Wolves and sheep move randomly though the space.
- Wolves and sheep die if they run out of energy.

¹<http://ccl.northwestern.edu/netlogo/models/WolfSheepPredation>

- Wolves eat sheep if they occupy the same space in the environment. Wolves gain *wolf-gain-from-food* units of energy from eating sheep. The sheep dies.
- Sheep eat grass if they are on a location of the environment that has grass. Sheep gain *sheep-gain-from-food* units of energy from eating grass. The grass dies in that grid location.
- Every time step, each sheep and each wolf has a chance (*sheep-reproduce* and *wolf-reproduce*) to reproduce asexually. Both the parent and the child split the parent's original energy evenly (i.e., parent's energy divided by two).
- Grass regrows after *grass-regrowth-time* number of time steps.




FIG. 1.3. The control and monitor interface for the Wolf Sheep Predation model.

The system-level concepts we are interested in are the number of sheep, the number of wolves and the number of grid locations containing grass. In NetLogo, these properties are displayed with monitors and a plot, as seen in Figure 1.3. The number of each population of agents may change continuously, but the average number of sheep converges. Another interesting feature is some ecosystems fail: either sheep or both sheep and wolves go extinct.




FIG. 1.4. Differences in populations based on changes of the *sheep-gain-from-food* parameter.

After working with this ABM for some time, a user will begin to realize that changes in the control parameters will yield different types of behavior. For example, by setting *sheep-gain-from-food* to 2, 3, and then 4, major differences in system-level behavior are apparent by viewing the graphs in Figure 1.4. When the value of *sheep-gain-from-food* is 4, the system rhythmically exhibits major changes in all three agent populations. When the

value is 2 or 3, the population remains relatively stable, but the average population values are different. When the value is low enough (e.g., 2) the wolves go extinct.

The Wolf Sheep Predation model is a good example of the intuitive disconnect between agent-level parameters and system-level properties. There is no clear *explicit* relationship between the controls presented in the user interface and the resulting system-level properties. An experienced user may have a qualitative understanding of the correlations, but would not be able to predict quantitative concepts, such as the average number of sheep after 2000 time steps. In Chapter 9: Results, I will show that the intuitive disconnect in this domain can easily be solved by AMF.

1.3 Overview of the ABM Meta-Modeling Framework

The foundation of this work is framing the problem of building a meta-model of an ABM as two sub-problems: the *forward-mapping problem* and the *reverse-mapping problem*. In Chapter 6: The Forward-Mapping Problem, I will discuss how AMF maps given values of the agent-level parameters to expected system-level property values with standard regression approaches. In Chapter 7: The Reverse-Mapping Problem, I will discuss how AMF maps a set of desired system-level property values to a set of agent-level parameters that would generate this behavior. My general approach to solving the reverse-mapping problem is to interpolate configurations using the forward mapping to approximate a smooth and continuous surface. This interpolated surface represents the space of configurations that would satisfy the system-level requirements set out by the user. Also, in Chapter 7, I will discuss alternative methods for solving the reverse-mapping problem.

AMF is simple and has only a few configuration points. This allows researchers to focus on the analysis of the system, instead of on the details of AMF. The framework consists of three major steps: sampling, solving the forward-mapping problem and solving

the reverse-mapping problem. In these three steps, the only configurations the user must perform are: define how to measure system-level properties of interest, provide the ranges of parameters to be sampled, and plug in a regression algorithm.

I will show in Chapter 9: Results that my framework is able to generate models of system-level behavior. For example, AMF is able to predict the number of sheep and wolves in the Wolf Sheep Predation model, given the configuration parameter values (the forward-mapping problem). Also, AMF is able to make suggestions for the values for the control parameters, given the the desired system-level property outcome (the reverse-mapping problem).

A more comprehensive overview of AMF is provided in Chapter 2: The ABM Meta-Modeling Framework.

1.4 Summary of Contributions

My main contribution presented in this dissertation is an in-depth analysis of meta-models of agent-based models. This analysis includes a discussion of methods for using regression to build models of the correlations between agent-level parameters and system-level properties. In addition, this dissertation contains a survey of ways that that meta-models can be used to inspect system-level behaviors of agent-based models.

The ABM Meta-Modeling Framework encapsulates my methodology for building meta-models of ABMs. The implementation of AMF as software serves as a proof-of-concept to show that my approach is implementable and applicable to a variety of domains. The software itself is a contribution, since it is available to be used by researchers interested in building meta-models of NetLogo ABMs. The design of the general framework is a contribution as well, since it could be implemented to interact with other agent-based modeling systems similar to NetLogo, or totally independent agent-based models.

1.5 Dissertation Organization

This dissertation is divided into nine chapters, including this one. Chapter 2: The ABM Meta-Modeling Framework explains each framework component in detail, explains how a new user would tailor AMF to a new ABM, discusses implementation details and gives an introduction to the forward- and reverse-mapping problems. Chapter 3: Related Work compares and contrasts approaches similar to AMF in motivation, with AMF. Chapter 4: Background provides information about NetLogo and algorithms used by AMF. Chapter 5: Defining System-Level Properties discusses the numerous different ways that system-level properties of ABM can be defined. Chapters 6 and 7 discuss my solutions to the forward- and reverse-mapping problems. Chapter 8: Using Meta-Models is a survey of different ways the meta-models generated by AMF can be used to analyze system-level properties of ABMs. Chapter 9: Results evaluates AMF on a domain-by-domain basis and provides explicit examples of how AMF has been used. Chapter 10: Conclusions and Future Work summarizes this dissertation, provides additional thoughts I have regarding this work and possible directions for future work.

Chapter 2

THE ABM META-MODELING FRAMEWORK

The ABM Meta-Modeling Framework builds meta-models that map the values of agent-level control parameter values to system-level property values, and vice versa. Learning these mappings are separate problems, which I call the *forward-mapping problem* and the *reverse-mapping problem*. To solve these problems, a user of AMF “plugs in” a regression algorithm of their choice. The framework uses this regression algorithm to learn the mappings and then provide the user with interfaces to query them. Once the mappings are learned, the user can query either for *prediction* or *control*. A prediction query uses the forward mapping to determine values for the expected values for system-level properties, given the system’s configuration. A control query uses the reverse mapping to suggest values for agent-level parameters, given desired values for system-level properties. Most of the implementation details and inner workings of AMF are abstracted away from the user, who only has to attend to a limited number of configuration points.

In this chapter, I will discuss the design goals of AMF, the framework structure, software implementation details, and how well AMF conforms to the design goals.

2.1 Design Goals of The ABM Meta-Modeling Framework

My specific goal in designing AMF was to make controlling and interacting with agent-based models more intuitive. In addition to this central goal, AMF strives to be:

- Domain independent – the design of AMF should minimize the amount of configuration for each new domain,
- Algorithm independent – any regression algorithm should be able to be plugged into AMF,
- Accurate – AMF should generate accurate predictions and control suggestions,
- Fast for the user – interactions with the models generated by AMF should require minimal computational time.

Domain independence is paramount because of the variety in which ABMs come. I have designed AMF such that the same general approach would work for any ABM. Also, I strove to minimize the amount of configuration needed to apply AMF to a new domain. These constraints I have set on the design make AMF broadly applicable to a number of domains, without the need of in-depth domain knowledge. To reinforce this claim, I have tested AMF on a number of diverse domains and used the same general approach for each.

Algorithm independence in a learning framework is important because different algorithms may be more effective for modeling different agent-based models. In general, the learning algorithms that will be discussed in this dissertation will satisfy the requirements for modeling most agent-based models. However, an in depth analysis of which types of algorithms should be used for different classes of ABMs is outside the scope of this dissertation research. In addition, algorithm independence allows AMF to scale with new advances in machine learning research, since future state-of-the-art regression algorithms can be plugged in just as easily as current approaches.

Accuracy and fast user response time appear to be obvious design goals. However, achieving these goals require sacrifices in performance in other portions of the framework. AMF requires a significant amount of computational time to sample different configurations of the target ABM. These large training sets can be used to build static meta-models of ABMs that are both accurate and fast to query. In contrast, an active learning approach would be able to learn models faster, but would require interaction with the user, increasing the amount interaction. Likewise, optimization approaches (e.g., hill climbing) could be used to generate arbitrarily accurate results, but typically require numerous iterations and would significantly increase the response time for a user’s query. This is because each step would have to run the ABM to calculate the fitness score, which could take several seconds. In summary, I am making the assumption that users studying ABMs with AMF are more interested in achieving more accurate results for their research and interacting with the models quickly, than spending less time sampling.

2.2 Framework Structure

The framework is split into several phases: sampling, solving the forward-mapping problem, solving the reverse-mapping problem, querying for prediction, and querying for control (i.e., suggesting a configuration). I explain with Figure 2.1 how the different phases interact with one another. Sampling makes observations from the actual ABM and then feeds the newly generated data set to the forward-mapping solver. The result of solving the forward mapping is a function f , which is used to predict behavior and to develop the reverse mapping f^{-1} . From an exterior perspective, the framework only has a limited amount of input and output: AMF takes in observations of an agent-based model and then provides interfaces to query for prediction and control. These queries interact with the user, as well as the agent-based model.

Many configuration points exist, which allow users of AMF to modify the behavior of the framework. However, the individual phases take the same output and provide the same output, no matter the configuration. This is what makes AMF a framework and not a collection of algorithms. For example, even though different regression algorithms can be used to learn the forward mapping, the forward mapping still provides predictions of how an ABM will behave, from the user's perspective.

2.2.1 Defining System-Level Behavior Properties

System-level properties of an ABM must be defined by the user of AMF. The measurement of a system-level property is a statistical or mathematical calculation based on the state of the ABM over some period of time (or “ticks”). The one assumption made by AMF is that the measurement is *stable*, meaning that as the same configuration is sampled repeatedly, the measurement's value should vary minimally. One way to measure stability is to calculate the standard deviation of measured values of several runs of the same system. The need for this assumption, ways to conform to it, and a more detailed explanation of how to define system-level properties is provided in Chapter 5: Defining System-Level Behaviors.

For example, there are several system-level behavior properties we measure in the Wolf Sheep Predation model. Three of the most obvious are the average number of sheep, average number of wolves, and average amount of grass over a significant number of ticks. These are measured by individually summing the number of sheep, wolves, and grass living at each time step and dividing by the number of ticks. Although the number of sheep and wolves change rhythmically, the average values typically converge to a single value after about 10,000 ticks.




FIG. 2.1. An overview of the phases of AMF and how data flows between them.

2.2.2 Sampling

The first phase is *sampling*. In this phase, numerous observations are made on different configurations of an agent-based model. A data set is generated that contains the independent variables (the agent-level control parameter values) and the resultant system-level behaviors that are measured, for each observation. This process can be significantly time consuming, dependent on several factors:

- Execution time of the model – the models will be executed numerous times, so the longer a model takes to execute, the longer the whole set of experiments will take to execute,
- Granularity – more detailed samplings will take more time, since more points must be sampled,
- Dimensionality – more agent-level control parameters result in a larger search space and naturally more points to sample,

A 120 observation sample of the NetLogo *Fires* model¹ (Wilensky 1997b), requires about thirty seconds² to generate. Meanwhile, a large sample of 50,000 observation of a Reynolds boid flock (Reynolds 1987) took approximately four days.

The user is required to have a minimal amount of domain knowledge to specify to AMF which ranges of values should be sampled. AMF is not able to automatically infer which value ranges are interesting, and thus must be explicitly defined.

This sampling process is easily parallelizable. Since each experiment is independent of one another, the set of all experiments can be segmented among a number of systems and processors to significantly reduce the computation time. Once all of the experiments are done executing, the results can be merged into one data set.

¹See Chapter 9 for detailed results

²Most experiments are executed on a 2.2GHz Intel Core 2 Duo running Mac OS X

For the scope of this dissertation, I limit AMF to use a simple random sampling method (i.e., randomly select points within a specified range) or a systematic sampling method (i.e., given ranges, sample evenly spaced points). I acknowledge that an intelligent sampling strategy can be improve the performance of many machine learning techniques, however, none are used for the sake of focusing on other portions of the framework. In future work, different sampling techniques' effect on accuracy and speed could be measured.

2.2.3 The Forward-Mapping Problem

Develop a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ that maps a provided configuration vector $\mathbf{x} = \{x_1, x_2, \dots, x_n\}$ to several system-level behavior properties $\hat{\mathbf{y}}$:

$$\hat{\mathbf{y}} \leftarrow f(\mathbf{x})$$

The set of values \mathbf{x} consists of all the agent-level parameters (independent variables) in the data set provided by the sample step. An individual mapping is learned for each system-level behavior measurement provided by the user.

This problem is solved with straightforward regression, such as k-nearest neighbor. In this phase of AMF, the regression algorithm is “initialized,” if necessary. For example, linear regression would need to solve the least-squares problem. Meanwhile, an algorithm like k-nearest neighbor would not need to initialize anything.

Once this phase is completed, the user is presented with f , an interface to the mapping built by the regression algorithm. The forward mapping is primarily used to *predict* what the system-level property values will be, given the system's configuration. For example, a learned mapping could be used to determine the average number of sheep given a configuration vector, without having to run the system. Some sample predictions, given particular system configurations, are shown in Table 2.1. The values for a system configuration rep-

resent *grass-regrowth-time*, *sheep-gain-from-food*, *wolf-gain-from-food*, *sheep-reproduce*, and *wolf-reproduce*, respectively.

Table 2.1. Sample Predictions of Behavior in the Wolf Sheep Predation Model

| Configuration | Average # Sheep | Average # Wolves | Average # Grass |
|-------------------|-----------------|------------------|-----------------|
| (30, 4, 20, 4, 5) | 162.8 | 76.1 | 964.8 |
| (30, 3, 26, 7, 5) | 122.8 | 90.4 | 1135.2 |
| (14, 3, 26, 7, 5) | 144.3 | 163.8 | 1621.4 |
| (5, 3, 17, 7, 5) | 946.9 | 3.4 | 1065.4 |

A more in depth definition of the forward-mapping problem, specific examples of regression methods used in AMF, and the role of regression is given in Chapter 6: The Forward-Mapping Problem.

2.2.4 The Reverse-Mapping Problem

Produce a mapping $f^{-1} : \mathbb{R} \rightarrow \hat{S}$ from a given system-level behavior properties \mathbf{y} to a set of configurations $\hat{S} = \{\hat{\mathbf{x}} | f(\hat{\mathbf{x}}) = \mathbf{y}\}$ (i.e., S is the set of behaviors that will produce behavior \mathbf{y}):

$$\hat{S} \leftarrow f^{-1}(\mathbf{y})$$

The problem of developing the mapping f^{-1} is what I call an “inverted regression problem” and has many unique challenges. The main challenge is f^{-1} does not describe a functional mapping. This is because f^{-1} describes a one-to-many relationship, since f is many-to-one. Therefore, AMF cannot use standard regression techniques to solve this problem. Instead, the default behavior of AMF is to approximate the inverse of the forward mapping, with a method that I developed. This approach has the benefit of using the forward mapping to develop the reverse mapping, so no additional input from the user

or data set are needed.

The reverse mapping can be used to suggest a system configuration that will exhibit a specific system-level properties. I call this process *control* of the ABM, since it is controlling the ABM at a system-level by suggesting configurations. For example, the reverse mapping could be used to suggest a configuration of (30, 4, 20, 4, 5) for a desired system-level property of 162.8 average sheep (see Table 2.1). However, using the reverse mapping is more involved than using the forward mapping, because f^{-1} returns a set of possible solutions, not a singular suggestion. If the mapping is to be used for control, any configuration in this set will satisfy the system-level requirements. Therefore, an additional step must be taken to extract a point from the set if it is to be used to control an ABM.

The implementation details of developing reverse mappings and how to use the reverse mappings are discussed in Chapter 7: The Reverse-Mapping Problem.

2.2.5 Summary of Configuration Points

The following is a summary of configuration points discussed in this section. For each new domain, the user must specify the following:

- A list of agent-level control parameters, and within what ranges they should be sampled,
- A list of system-level behavior properties and the way to measure them.

In addition, the user may configure the following if they wish to stray from the defaults:

- The sampling strategy (default: random sampling),
- The regression algorithm to be used for the forward mapping (default: k-nearest neighbor),

- The method for the reverse mapping (default: approximate the inverse of the forward mapping).

2.3 Software Implementation Details

So far, I have discussed AMF abstractly and have avoided specific implementation details. This is because the framework is a methodology and could be reimplemented in a number of different ways or to work with a number of different ABM simulation systems. In this section, I discuss the details of my proof-of-concept implementation that was used to run many of the experiments documented Chapter 9: Results.

My implementation works directly with NetLogo³ to sample, predict and control NetLogo ABMs. All interactions with NetLogo, such as running experiments, retrieving a current configuration or pushing a suggested configuration are handled with Java through NetLogo's Java API. The rest of the framework is implemented in Python. Each step of AMF returns its result as a file, so that it can be passed to the next step or saved for later use. For example, the forward mapping writes the model to a file so that it can be used by both the prediction script and the reverse mapping script.

Since each domain has different variable names and nuances, the user must implement a Java program that interacts with NetLogo's API. An abstract base class for interaction is provided to guide the development of this program. Next, the user lists the configuration parameters and the ranges of these values so that the sampling can begin. Measuring the system-level parameters can either be calculated within NetLogo or in Java. For example, I modified the standard Wolf Sheep Predation model to keep track of the number of sheep at each tick. Then, to extract the value, NetLogo's API is used to retrieve the sum of the sheep divided by the number of ticks. A similar calculation could be performed within

³More about NetLogo is covered in Chapter 4: Background.

Java program by retrieving each piece individually, then performing statistics on them. Performing the statistics in Java have the benefit of not having to change the NetLogo ABM source code. The data set is written to a file, row by row. Therefore, the results of several instances of the sampling, executing on different machines, can be easily concatenated. An example of a sampling program for Wolf Sheep predation is given in Appendix A.1.

The regression algorithms used with AMF must conform to a standard API⁴ and are passed as arguments into the forward, reverse mapping and prediction scripts. Therefore, no configuration of these core scripts is needed, since they are “pluggable.”

An overarching “master script,” written in Python, runs all steps automatically, limiting the amount of direct interaction with the framework software.

In addition to the core framework software, I have developed a toolkit that queries the models generated by the forward- and reverse-mapping models. The core tools include:

- Query the forward mapping for a prediction by passing in a configuration,
- Query the reverse mapping for a set of possible solutions by passing in a desired system-level behavior configuration,
- Visualize the mappings.

Each of these will be discussed in detail in Chapter 8: Using Meta Models.

2.4 Analysis of AMF vs. the Design Goals

In summary, I align the actual implementation of AMF with the design goals.

Domain independence— My framework’s implementation is domain independent, to an extent. AMF reduces the forward-mapping problem to a classical regression problem

⁴The specification for the standard regression algorithm interface is given in Appendix B.

of learning the correlation between the agent-level parameters and the system-level properties. My reverse-mapping problem solution is also domain independent because it interacts exclusively with the forward mapping, which is domain independent.. However, my implementation of AMF still requires the user to specify the agent-level parameters and how to measure the system-level properties. This is a reasonable, since automatically detecting configuration points and having a computer determine behaviors of interest in an ABM would be a challenging unsupervised learning problem.

Algorithm independence— Any regression algorithm can be plugged into my implementation, as long as they conform to the standard API. A simple Python wrapper can be written to adapt an already existing third party regression algorithm to work with my implementation of AMF. My default process of inverted regression requires nothing special of the regression algorithm, because the inversion learning process uses the algorithm's standard forward-mapping behavior.

Accuracy— No extra error is incurred by AMF, itself; the accuracy of AMF depends on the regression algorithm used and the amount of time spent sampling. If AMF is not providing accurate results with state of the art regression algorithms, either the correlations cannot be learned with current technology or not enough time has been spent sampling.

Fast for the user— Most computation time is spent sampling and learning the models. These operations are offline and do not affect the response time of real-time interaction with the user. The response time of the forward mapping and the reverse mapping are proportional to the running time of the regression algorithms, but are often require less than a few seconds.

Chapter 3

RELATED WORK

3.1 Experimentation in ABMs

Experimental platform for messing around with ABMs: (Bourjot – “A platform for the analysis of artificial self-organized systems” 2004) (relevant?)

ABMs have been used to study behaviors in biological systems... Domain specific work interested in system-level behavior: ant lane formation (Couzin & Franks 2003), marching locusts (Buhl *et al.* 2006), fish schools (Parrish, Viscido, & Grunbaum 2002). Most of these studies are qualitative in nature.

particle swarm optimization: empirical study (Shi & Eberhart 1998); more general approach: (Van den Bergh & Engelbrecht 2006)

3.2 Prediction of System-Level Behavior

physics-based control policy (Spears *et al.* 2004) – similar to our approach, but the models are tightly coupled to the domain. the system was designed with system-level models in mind. Algebraic inversion for inverse mapping is nice. Inspires the nonlinear regression approach in AMF.

Macroscopic models of swarm robot systems (Lerman & Galstyan 2002)(Lerman, Martinoli, & Galstyan 2005) – similar in motivation, but models the system in a more

specific way (FSAs). Specific to systems in which agents can be modeled as FSAs. Our approach is more general, since it just looks at parameters.

3.3 Inversion of Neural Networks

Inversion of neural networks:

- (A Linden and J Kindermann “Inversion of multilayer nets” 1989 – optimization problem solved by gradient descent)
- (Bao-Liang Lu “Inverting Feedforward Neural Networks using Linear and Nonlinear programming” 1990 – formulate the inverse problem as a nonlinear programming problem, a separable programming problem or a linear programming problem)
- (S. Lee and R.M. Kill “Inverse Mapping of continuous functions using local and global information” 1989 – iterative update towards a good solution)
- (Michael I. Jordan work with robot arm “Forward Models: supervised learning with a distal teacher” – asks the question, what configuration of the robot arm will yield this behavior?)

All of these are optimization techniques. They do not return an actual mapping. Also, some of the techniques are restricted to neural networks (and are thus not algorithm-independent).

Chapter 4

BACKGROUND

The purpose of this chapter is to provide necessary background information for understanding concepts related to AMF. In contrast to Chapter 3: Related Work, the previous research presented in this chapter do not share my motivation in AMF, but are a foundation for which AMF is built upon.

The ABM Meta-Modeling Framework uses preexisting research in two major areas: agent-based modeling and regression. The AMF is targeted at predicting and controlling behavior in an agent-based model (ABM). The next section discusses what an ABM is, discusses some examples of ABMs, and lists some existing multi-agent software frameworks. In the following section, I define regression, outline favorable and unfavorable properties that regression algorithms, and outline a number of regression algorithms that have been used by AMF. Finally, I discuss methods that are used in solving the reverse-mapping problem: plane intersection and optimization.

4.1 Agent-Based Modeling

Agent-based models are implemented from the agent perspective. Agents in ABMS are typically:

- bounded by a limited global view,

- perform local interactions, affecting their local environment and neighboring agents,
- are autonomous (i.e., not following some top-down control), and
- may be heterogeneous (i.e., agents within the system can have different properties).

(Epstein 1999) This is opposed to an “observer” perspective, in which an overarching control system dictates what agents are to do. For example, in NetLogo, an agent moves forward two units with the following command:

```
> fd 2
```

Notice that this code is agnostic to the global direction of the agent or the position of the agent in the environment. The code simply tells the agent to move forward two units. In contrast, if this were to be done from the observer perspective in NetLogo, the following code would be necessary:

```
> set [xcor] of turtle 0
([xcor] of turtle 0 + 2 * cos([heading] of turtle 0))
> set [ycor] of turtle 0
([ycor] of turtle 0 + 2 * sin([heading] of turtle 0))
```

This code computes the change in the x-coordinate and the y-coordinate, given that the agent should move 2 units. Then, it adds the result to the agent’s original x- and y-coordinates. Finally, it sets this turtle’s x- and y-coordinates to the new x- and y-coordinates. This process is cumbersome. The agent-based property allows for a more intuitive way to build multi-agent systems.

One of the common uses of ABMs is to discover which local interactions generate a given emergent behavior of a system, through experimentation. The research question is posed well by Epstein as the *Generativist’s Question*:

How could the decentralized local interactions of heterogeneous autonomous agents generate the given regularity?

To answer this question, Epstein then poses the *Generativist's Experiment*:

Situate an initial population of autonomous heterogeneous agents in a relevant spatial environment; allow them to interact according to simple local rules, and thereby generate—or “grow”—the macroscopic regularity from the bottom up. (Epstein 1999)

If a model accurately generates the emergent behavior of the target system, then that ABM could explain why that behavior emerges. Agent-based models are particularly well-suited to answer the Generativist's question.

The concept of ABMs can be extended to other uses beyond modeling. Many swarm intelligence techniques, such as particle swarm optimization (Kennedy & Eberhart 1995) and ant colony optimization (Dorigo 2004), use a decentralized agent-based approach to solve optimization problems. Although they are inspired by naturally occurring phenomena, their purpose is entirely separate.

Implementing systems as an ABM could be more natural than other approaches when the system cannot be defined in an aggregate manner or when individual behavior is complex (Bonabeau 2002). ABMs are often compared to equation-based modeling (EBM), in which the model is a parametrized system of equations that describe system-level behavior (i.e., the observer perspective). For example, a supply chain management system could be modeled as an ABM or an EBM (Parunak, Savit, & Riolo 1998). In the ABM proposed by Parunak et. al, individual agents represent different companies that trade with one another. In contrast, the EBM is a series of ordinary differential equations describing the input and output of different components of the supply network. It is difficult to model certain behavior, such as changes in state, erratic behavior and local interactions

with these observer-perspective equations. In my example above where the agent moves forward, a relatively simple task of adjusting the position of an agent proved to be overly complicated. When more detailed operations need to be performed, the complexity of the observer-perspective model increases drastically. Agent-based models remedy this by simply changing the context of the programming.

4.1.1 Examples of Agent-Based Models

ABMs have been used to model a variety of different systems. In this subsection, I outline a number of agent-based models that have been developed by other researchers in the past.




FIG. 4.1. A boid flock moving through a two-dimensional space.

Reynolds Boid Flocking One of the first popular ABMs was the boid flock (Reynolds 1987)(Reynolds 1999), in which agents move through an environment in a similar way to how birds flock. A screen shot from NetLogo's boid implementation (Wilensky

1997a) is shown in Figure 4.1. Agents follow three simple rules:

- move away from other agents to avoid collisions,
- align to move in the same direction as nearby agents, and
- move towards the central position of local flockmates.

These simple agent-level behaviors result in the elegant flocking behavior observed from a top-down view of the system. This system naturally lends itself to being implemented as an agent-based model because the agents act autonomously and have local interactions with only their neighbors. The flocking behavior is emergent and therefore would be more difficult to program directly.

Boid flocking has been used in a number of applications, such as visualizing time-varying data (Moere 2004), clustering documents (Cui, Gao, & Potok 2006), controlling unmanned air vehicles (Crowther & Riviere 2003) and art (Boyd, Hushlak, & Jacob 2004).

Social Insect Behavior A number of different projects have aimed at modeling insect societies. In modeling these systems, researchers are able to gain insights into how individual agent behaviors affect their emergent behavior. These models are typically based on observations of the real insects.

An agent-based model approach has been used to model how army ants form traffic lanes (Couzin & Franks 2003). The particular species of army ants discussed in this work form two-way highways to reduce the amount of head-on collisions while ants are searching for food and returning food. The authors show how the turning rates and perception of the individual ants affect the efficiency of the ants. They found that when the turning rate is too high, ants are too willing to steer off course and intersect the path of other ants. On the other hand, when the turning rate is too low, ants will not adjust their heading to avoid head on collisions. By tuning their model to optimize the highway performance, they reached

an accurate model of the traffic flow in army ants. The result of this research is an accurate model that describes the individual behavior of ants in a traffic situation.

With a different species of ants, researchers were able to produce an agent-based model that simulates the process of collectively selecting a location for a nest (Pratt *et al.* 2005). This process is interesting because an entire ant colony converges on one location, even though many of the ants only have scouted one nest site. Agents are implemented as state machines, in which agents are either exploring for a site, assessing a site, canvassing a site or committed to a site. In each of these phases, the agent performs different actions and at any time may reject the current site and begin exploring once again. Over time, all ants converge on a single site as a nest. With this model, the authors showed that a colony-level decision can be made by ants following agent-based rules.

An ABM has been developed to simulate swarming locusts (Buhl *et al.* 2006). Locusts have an interesting property in that when they are isolated, they tend to not stray from their current location. However, when locusts are accompanied by several other locusts, they begin to “march,” travelling from one area to another, consuming everything in their path. The authors were particularly interested in determining at what density the locusts will begin to march. By changing the number of simulated locusts in a confined space, they were able to determine at what critical density they would march. This simulation approach is far more practical than experimenting with real locusts. The authors suggest that their models could be used to predict when locust swarms will occur to help warn farmers to protect their crops.

4.1.2 Models of Human Societies

Although human behavior is more complicated than insect behavior, local human interactions can often be generalized to build accurate models of a subset of human society. Most of these aim to create an accurate as possible model of the human interactions with

others in order to predict some society-level behavior.

Agent-based models have been used to determine how locals would react to an incident at the Pacific Missile Range Facility (PMRF) in Hawaii (Zanbaka, HandUber, & Saunders-Newton 2009). The researchers used census data to represent each islander as an individual agent that either has a positive sentiment or a negative sentiment towards the missile facility. Agents can either change its sentiment towards PMRF by either interacting with another agent that has a different sentiment or experiencing a local event. The model simulates local events and agent interactions and displays how either positive or negative sentiment propagates throughout the island. With this model they were able to model hypothetical situations and how they effect the sentiment on the island.

EpiSims is an agent-based simulation tool that models disease outbreaks (Eubank *et al.* 2004). It uses estimates of how diseases are transmitted and how humans interact based on census and land-use data to realistically simulate the human society. The system simulates individuals going to work and shopping, which exposes them to the disease and exposes the disease to others. With this system, researchers are able to predict the effectiveness of mass and targeted vaccination strategies, given a particular community.

STREETS is an agent-based model of pedestrian traffic (Schelhorn *et al.* 1999). Pedestrian traffic is affected by two major aspects: the layout of the street network and the location of attractions. Instead of trying to analyze the behavior of the system from this data alone, STREETS simulates people to determine a number of properties of pedestrian traffic patterns. STREETS initializes the system with a statistically accurate distribution of individuals across the environment. Next, the agent-based model simulates the movement of agents from their arrival point to and through the urban center. Agents visit buildings and attractions and walk between them. Researchers use this tool to observe a top-level view that can be used to analyze the effectiveness of the urban center layout. Hypothetical modifications to the environment can be performed in simulation to determine if it would

improve the traffic situation.

4.1.3 Multi-Agent Software Frameworks

Implementations of different multi-agent systems have many similarities. To reduce the amount of “boiler plate” code for each new multi-agent system, a number of unique multi-agent software frameworks have been developed in the past two decades.

The Swarm Simulation System Swarm is one of the original agent-based modeling platforms, originally developed at the Santa Fe Institute in the 1990s. Swarm is currently supported and hosted¹ by an independent organization called the Swarm Development Group. The developers’ motivation was to enable researchers to focus less on implementation and more on actual experimentation (Minar *et al.* 1996). To achieve this goal, they implemented a number of libraries in Objective-C and Java that aid in the creation of “swarm” objects, the building block of a Swarm simulation. The swarm objects manage the agents and handle the time schedulers that queue the order of agent actions. Also, agents can be nested as hierarchies in swarm objects, to allow for different types of agents. Swarm is a discrete time simulation, which means time progresses as actions occur.

Programming in Swarm can require a significant amount of programming overhead, since the libraries are not a single integrated application (Klein 2003).

Repast The Recursive Porous Agent Simulation Toolkit (Repast) is a comprehensive agent-based modeling toolkit, focusing on modeling social interactions (Collier 2003). The toolkit is freely available to download.² Repast is fully object oriented and attempts to be as platform independent as possible, supporting programming in Java, Python, Visual

¹Download Swarm at <http://www.swarm.org/>

²Download Repast at <http://repast.sourceforge.net/>

Basic.Net, and more.

Repast's features are split into six modules (North, Collier, & Vos 2006):

1. The engine module – the core module that controls the agents, environment and scheduler,
2. The logging module – records execution results,
3. The interactive run module – manages user interaction with the model,
4. the batch run module – allows the user to set up a set of simulations to be ran in succession,
5. the adaptive behaviors module – an optional module that provides built-in adaptive agent behaviors, that use techniques such as genetic algorithms and neural networks,
6. the domains module – an optional module that helps define environments, such as social systems, geographic information systems, and computational game theory.

Repast's feature set is quite comprehensive and extensible.

Breve Breve is a unique 3D simulation environment that focuses on decentralized systems and artificial life (Klein 2003). Like most other simulation environments, Breve is freely available.³ The authors of Breve tout that it aims to provide a platform for physically realistic 3D models. The models simulate continuous time, and continuous space, unlike other toolkits which model time as discrete events and have grid world environments. In addition, some physics such as gravity and object collision resolution with friction are built-in features of every simulation. Breve's OpenGL display engine allow for easy to

³Download Breve at <http://www.spiderland.org/>

implement 3D-accelerated graphics. Users of Breve must use a custom object-oriented language called Steve.

One of the motivating applications for building Breve was Sims' (1994) evolved 3D creatures (Klein 2003). In this project, creatures composed of several blocks and joints compete in a game to be closest to a ball. The features of Breve naturally fit to this domain, since it has physically realistic servos and objects that can interact with one another. Due to these features, Breve is an interesting system for modeling individual mobile agents, as well as multiple mobile agents.

MASON MASON is a multi-agent simulation toolkit developed at George Mason University and is freely available.⁴ MASON takes a different approach than previous multi-agent toolkits in that it strives to be minimalist and efficient for up to a million agents (Luke *et al.* 2005). It is meant to be ran on a number of back-end computation servers in parallel, without visualization. MASON does not have any domain-dependent tools or built-in environments like Breve or Repast, leaving most to the user to implement in Java. Although MASON is quite minimalist, it is designed to be extended so that it can be used for either simple simulations or as a foundation for new simulation systems. MASON simulations can be interacted with a separate visualization that binds to the simulation. This is a different paradigm than other toolkits like NetLogo that are tightly coupled with the visualization.

NetLogo NetLogo is a relatively new ABM system, which was started in 1999 out of Northwestern University as a derivative of StarLogo (Tisue & Wilensky 2004). The software is free to download,⁵ but not open source. As the name suggests, the language

⁴Download MASON at <http://cs.gmu.edu/~eclab/projects/mason/>

⁵Download NetLogo at <http://ccl.northwestern.edu/netlogo/>

used by NetLogo is a derivative of Logo, a Lisp-like language. An artifact of this language that remained in NetLogo is that agents are referred to as “turtles.”

NetLogo’s language is tailored to the agent-based model paradigm. Particularly, a program can change context to an individual and execute code from its point of view. Also, it can execute these individual actions concurrently to a population of agents. For example, the code to ask all the agents to move forward one unit, then turn right would be:

```
ask turtles [ fd 1 rt 90 ]
```

In addition to turtle agents, the other basic agent type is the “patch.” Patches are organized in a grid in the two-dimensional environment that agents are confined to. Users can interact with patches from their context, like the turtles. Also, it is easy to retrieve turtles that are in contact with a patch and to retrieve which patch a turtle is on, from their respective contexts. This makes turtle interactions with the environment simple to implement.

NetLogo also includes a built-in user interface and user interface editor. The interface consists of controls, monitors and the domain visualization. Users define controls which bind to global variables and buttons that bind to function calls. In addition, users can add monitors and plots, which show a variable’s value or plot a variable’s value over time. These are useful tools in conveying information that the domain visualization cannot. NetLogo can be run “headless,” to facilitate parallel execution of models or to reduce run-time.

NetLogo comes with a number of extensions, most notably BehaviorSpace and HubNet. BehaviorSpace is a tool for setting up a set of systematic experiments with different system configuration parameters. BehaviorSpace performs a user-defined measurement value for each run and reports the result to a data set in the form of a spreadsheet. HubNet is a server/client tool that allows several users of a system to interact with a NetLogo ABM at the same time, which is useful for classroom instruction.

NetLogo can be interacted with from an external Java API. A NetLogo “workspace”

is instantiated as an object and can be sent NetLogo commands and can be queried for a current value. This is useful for an number of tasks. First, sometimes NetLogo's language is not expressive enough and some may find programming in Java more familiar. In addition, external Java libraries (such as machine learning libraries) can be used without modification by integrating with NetLogo from its Java API. Also, a sequential experiment system can be implemented in Java if a user would like to run experiments without BehaviorSpace.

An extensive model library with over 140 sample models are bundled with the NetLogo distribution. These already existing models serve as excellent examples for new models as well as starting points for modified models. Most of the domains discussed later in this dissertation are from this model library.

NetLogo is my ABM system of choice for this dissertation for a number of reasons. First, NetLogo's learning curve is surprisingly short. From personal experience, a new user can learn to write a domain similar to the Wolf Sheep Predation model in less than a couple hours. In addition, the documentation⁶ is well organized and detailed. Second, the model library provided several models that are interesting and easy to work with. Since each of these models are implemented in NetLogo, I was able to implement a system that interacted with each of these in a similar way. Other multi-agent system toolkits provide more flexibility in implementation, which would make identifying agent-level control parameters and system-level properties more difficult. Third, although NetLogo's platform is well contained, the interactions possible with the Java API are superb. I am able to implement my framework as an external modular application that interfaces with NetLogo, instead of a built-in system dependent tool. This will make my research easily extensible to other ABM systems in the future.

⁶The NetLogo documentation is available at <http://ccl.northwestern.edu/netlogo/docs/>

4.2 Regression

Regression is an integral part of this dissertation research, as it is the foundation of solving the forward- and reverse-mapping problems. Regression techniques are used to predict values of dependent variables, given values for the independent variables. The typical approach is to use a model developed from a sample training set to infer new values that have not been sampled. Also, regression can be used to smooth the natural error in sampling.

Linear regression is perhaps the simplest form of regression that fits a line to fit one independent variable with one dependent variable. It takes the form of:

$$y = mx + b.$$

There exists a closed-form solution to determine m and b in order to minimize the sum of least-squares between the training data and the curve.

Unfortunately, linear regression is limited in that it only models linear relationships, and is not able to model nonlinear ones, such as rhythmic oscillations, or quadratic correlations. Two paths are generally taken to address this concern. The first is to use the kernel trick (Muller *et al.* 2001) to map nonlinear data to be linear, but higher dimensional. Then, once the data is projected into a space in which it is linear, linear methods (e.g., linear regression, perceptrons (Minsky & Papert), or support vector regression (Smola & Schölkopf 2004)) can be used.

The second path are nonlinear methods, which relax the problem to learning a function $f(\cdot)$ that satisfies the following:

$$y_i = f(\mathbf{x}_i) + \varepsilon_i.$$

In this equation, y_i is a dependent variable, \mathbf{x}_i is the vector of independent variables asso-

ciated with y_i . Also, there is an assumption of some normally distributed and independent error with each point, denoted by ε_i . The process of *nonparametric* regression is to estimate this function directly (Fox 2002).

Alternatively, a *parametric* regression model is one in which the form of the function $f(\cdot)$ is relatively fixed. Instead of learning $f(\cdot)$ directly, a vector of model parameters β are optimized to fit the data (Fox 2002):

$$y_i = f(\beta, \mathbf{x}) + \varepsilon_i.$$

For example, $f(\cdot)$ could model two-dimensional sinusoidal data with the model:

$$f(\beta, (x_1, x_2)) = \beta_1 + \beta_2 \sin(x_1 - \beta_3) + \beta_4 \sin(x_2 - \beta_5)$$

Typically, the value of the sum of least-squares metric is minimized to fit the curve to the data.

While developing the ABM Meta-Modeling Framework, I needed regression algorithms that satisfied a number of properties. The regression algorithm should scale well with dimensionality, since agent-based model behavior spaces (i.e., the space of agent-level parameters and system-level properties) are typically highly dimensional. I also needed algorithms that could model nonlinear correlations, since behavior spaces are expansive and rarely maintain linearity throughout the entire space. However, in the ABMs I have experimented with, all spaces are very smooth locally. That is, there is little variance in values between nearby locations in the behavior space. Even in the situation where there is variance, the error is typically evenly distributed and cancels out if the same configuration is sampled a number of times. Therefore, nonlinear parametric regression can fit to the overall behavior space, without having to worry about fine grained variations in the data.

Also, the mappings build by the regression models should be smooth. This is important for the tractability of finding the reverse mapping, since bumpy surfaces will yield more intersections than smooth ones. Therefore, regression methods that smooth the data, such as LOESS, are effective in this situation.

[TODO: an image showing graphs of the fires, boids, and sheep/wolf predation behavior spaces (need to make wolf/sheep data set to do this)]

In the course of testing AMF, I used four regression techniques: k-nearest neighbor, LOESS, least-squares nonlinear regression, and Multilinear Interpolation.

4.2.1 K-Nearest Neighbor Regression

K-nearest neighbor (KNN), in general is the process of selecting the k nearest points to some point x . These points can then be used for classification or regression. In the simplest case, each neighbor's y value is averaged to inter a value y' for the given point x . KNN is easy to implement and accurate with large data sets. Accuracy of this regression can be improved by weighting points by distance or by using a secondary regression algorithm (e.g., linear regression) on the neighbors.

We use KNN in our results as a baseline, since it is easy to implemente and accurate with large data sets. In addition, it is nonparametric, which makes it agnostic the the shape of the behavior space. This results in KNN requiring no configuration or tailoring to a particular domain.

However, when implemented naively, the run-time scales linearly with the size of the data set because the process of finding the nearest neighbors requires a scan of all data points. This search time can be reduced with the implementation of faster searching techniques, such as locality sensitive hashing (Gionis, Indyk, & Motwani 1999), among other methods. Even with faster neighbor locating techniques, some sort of search is required for each query to KNN.

The spaces generated by KNN are not continuous (particularly if the neighbors are not weighted). This is because similar points will have the same neighbors and, in effect, have the same predicted value. There are separations in which on one side the k -nearest neighbors will include a particular data point, meanwhile on the other side it will not. At this separation, a noncontinuous jump will occur in predicted values. This can be a problem for solving the reverse-mapping problem because the intersection between hyperplanes generated by KNN will not intersect smoothly. This can be remedied by using KNN in conjunction with multilinear interpolation.

4.2.2 Robust Locally Weighted Regression and Smoothing Scatterplots (LOESS)

Robust Locally Weighted Regression and Smoothing Scatterplots (LOESS) is a mature nonparametric regression technique for fitting curves to data (Cleveland 1979)(Cleveland & Devlin 1988). Similar to k -nearest neighbor, LOESS uses local points to infer values of data points. LOESS also iteratively smooths the regression curve, allowing for variable magnitudes of smoothing. This is important for particularly noisy data sets. The role of LOESS in the testing of AMF is it is more robust than KNN, yet still nonparametric.

Abstractly, LOESS fits a new \hat{y}_i to every \mathbf{x}_i to replace the original noisy y_i . Then, LOESS iteratively smooths each \hat{y}_i to reduce variability. Specifically, LOESS follows the following steps:

1. Initialization step – for each \mathbf{x}_i , weight all other points x_j , where $j = 1 \dots n$, based on a weighting function $W(\mathbf{x}_i, \mathbf{x}_j)$. W weights closer points higher. Perform weighted linear regression to infer \hat{y}_i .
2. Smoothing step – For each \mathbf{x}_i , weight all other points based on w_i , as well as a new weight function δ , which returns magnitude inversely proportional to $y_i - \hat{y}_i$ (i.e., the larger the distance, the lesser the weight). Perform weighted linear regression to

infer \hat{y}'_i .

3. Set each \hat{y}_i to \hat{y}'_i .

4. Repeat the the smoothing step as many times as necessary.

Values in between xs can be now be interpolated.

LOESS can be customized in a number of ways. The different weighting functions can be changed to change the amount of influence other points give to the linear regression step. Typically, these weighting functions have a “smoothing” parameter, when increases the influence of further away points, effectively reducing the strong influence of local points. Adjusting the smoothing parameter is important because if it is too low, the regression curve will be erratic, meanwhile, if the smoothing parameter is too high, some features of the curve may be eliminated.

LOESS is relatively computationally expensive in comparison to other approaches. This is because more weights are applied and points are readjusted iteratively a number of times. However, the model can be built offline and then stored. Once the \hat{y}_i values are inferred, they can be used to interpolate new points (similar to multilinear regression) in response to queries. This provides a quick response time to users wanting to infer to \hat{y} values. This fits one of the design criteria of AMF: the regression algorithm should be fast for the user.

The *R* statistical package⁷ has a built in LOESS function.

4.2.3 Nonlinear Regression

(Gallant 1975)

⁷R website: <http://www.r-project.org/>

4.2.4 Multilinear Interpolation

Multilinear interpolation is a dynamic approach that uses multi-dimensional interpolation between “knots” to generate a smooth surface across a space of any dimension (Davies 1997). Knots are sampled data points, scattered across the behavior space in a regular fashion such that the knots, when connected, form hypercubes. When a point \mathbf{x} is queried, multidimensional interpolation is performed using the corners of the hypercube to infer the value of \hat{y} .

The major downside to multilinear interpolation is that the sampling needs to be systematic and even spaced. To remedy this situation, another regression algorithm is used to compile a set of evenly spaced points. These evenly spaced inferred points are then passed to a traditional multilinear interpolation approach.

In this dissertation, I typically use multilinear interpolation as a supplement to the other regression approaches. A number of points can be inferred with another regression algorithm and then interpolated between with this approach.

4.3 Inverting Regression

4.3.1 Plane Intersection

(Patrikalakis *et al.* 1993)

4.3.2 Optimization

Chapter 5

DEFINING SYSTEM-LEVEL PROPERTIES

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

For example, in the Wolf Sheep Predation domain, one simple system-level measurement could be the average number of wolves over a thousand time steps. This number can be misleading, because certain configurations will sometimes result in the wolves going extinct (i.e., zero wolves). Other times, the same configurations will have the wolves converge to a stable non-zero population. Therefore, the expected value for the average number of wolves \hat{w} to be:

$$E(\bar{w}) = P(\textit{extinct}) * (\bar{w}|\textit{extinct}) + (1 - P(\textit{extinct})) * (\bar{w}|\neg\textit{extinct})$$

where *extinct* is whether the wolves went extinct or not, $P(\textit{extinct})$ is the probability wolves go extinct, and $(\bar{w}|\neg\textit{extinct})$ is the average number of wolves, given they did not go extinct. \hat{w} is not very stable, because sometimes it is zero and sometimes it is $(\bar{w}|\neg\textit{extinct})$.

Making a prediction based on the expected value of \hat{w} will always have a specific amount of error associated with it. To remedy this problem

Chapter 6

THE FORWARD MAPPING PROBLEM

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Chapter 7

THE REVERSE MAPPING PROBLEM

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Chapter 8

USING META-MODELS

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Chapter 9

RESULTS

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Chapter 10

CONCLUSIONS AND FUTURE WORK

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Appendix A

CODE SAMPLES

A bunch of code samples will be put here. These will serve as examples to show how to do things or to show how easy they are

A.1 Sample Java/NetLogo Sampling Program

This code example will show how to write a sampler for AMF.

Appendix B

INTERFACE SPECIFICATION FOR REGRESSION ALGORITHMS

I will explicitly detail the interface that a regression algorithm to be plugged into AMF needs to have.

REFERENCES

- [1] Bonabeau, E. 2002. Agent-based modeling: Methods and techniques for simulating human systems. *Proceedings of the National Academy of Sciences of the United States of America* 99(Suppl 3):7280.
- [2] Boyd, J.; Hushlak, G.; and Jacob, C. 2004. SwarmArt: Interactive art from swarm intelligence. In *Proceedings of the 12th Annual ACM International Conference on Multimedia*, 628–635.
- [3] Buhl, J.; Sumpter, D.; Couzin, I.; Hale, J.; Despland, E.; Miller, E.; and Simpson, S. 2006. From Disorder to Order in Marching Locusts. *Science* 312(5778):1402–1406.
- [4] Cleveland, W., and Devlin, S. 1988. Locally weighted regression: an approach to regression analysis by local fitting. *Journal of the American Statistical Association* 83(403):596–610.
- [5] Cleveland, W. 1979. Robust locally weighted regression and smoothing scatterplots. *Journal of the American Statistical Association* 829–836.
- [6] Collier, N. 2003. Repast: An extensible framework for agent simulation. *The University of Chicago's Social Science Research*.
- [7] Couzin, I., and Franks, N. 2003. Self-organized lane formation and optimized traffic flow in army ants. In *Proceedings of the Royal Society of London, Series B*, volume 270, 139–146.
- [8] Crowther, B., and Riviere, X. 2003. Flocking of autonomous unmanned air vehicles. *Aeronautical Journal* 107(1068):99–109.

- [9] Cui, X.; Gao, J.; and Potok, T. 2006. A flocking based algorithm for document clustering analysis. *Journal of systems architecture* 52(8-9):505–515.
- [10] Davies, S. 1997. Multidimensional triangulation and interpolation for reinforcement learning. *Advances in neural information processing systems* 1005–1011.
- [11] Dorigo, M. 2004. *Ant Colony Optimization*. MIT Press.
- [12] Epstein, J. 1999. Agent-based computational models and generative social science. *Complexity* 4(5):41–60.
- [13] Eubank, S.; Guclu, H.; Anil Kumar, V.; Marathe, M.; Srinivasan, A.; Toroczkai, Z.; and Wang, N. 2004. Modelling disease outbreaks in realistic urban social networks. *Nature* 429(6988):180–184.
- [14] Fox, J. 2002. *An R and S-Plus companion to applied regression (Appendix)*. Sage Publications.
- [15] Gallant, A. 1975. Nonlinear regression. *The American Statistician* 29(2):73–81.
- [16] Gionis, A.; Indyk, P.; and Motwani, R. 1999. Similarity search in high dimensions via hashing. In *Proceedings of the 25th International Conference on Very Large Data Bases*.
- [17] Kennedy, J., and Eberhart, R. 1995. Particle swarm optimization. In *Proceedings of the 1995 IEEE International Conference on Neural Networks*.
- [18] Klein, J. 2003. Breve: a 3D environment for the simulation of decentralized systems and artificial life. In *Proceedings of the Eighth International Conference on Artificial Life*, 329–334. MIT Press.

- [19] Lerman, K., and Galstyan, A. 2002. Mathematical model of foraging in a group of robots: effect of interference. *Autonomous Robots* 13(2):127–141.
- [20] Lerman, K.; Martinoli, A.; and Galstyan, A. 2005. A review of probabilistic macroscopic models for swarm robotic systems. In *Swarm Robotics Workshop: State-of-the-art Survey*, 143–152. Springer.
- [21] Luke, S.; Cioffi-Revilla, C.; Panait, L.; Sullivan, K.; and Balan, G. 2005. Mason: A multiagent simulation environment. *Simulation* 81(7):517–527.
- [22] Minar, N.; Burkhart, R.; Langton, C.; and Askenazi, M. 1996. The swarm simulation system: A toolkit for building multi-agent simulations.
- [23] Minsky, M., and Papert, S. Perceptrons: expanded edition. *MIT Press, Cambridge MA* 19882:20.
- [24] Moere, A. 2004. Time-varying data visualization using information flocking boids. In *Proceedings of the 2004 IEEE Symposium on Information Visualization*, 97–104.
- [25] Muller, K.; Mika, S.; Ratsch, G.; Tsuda, K.; and Scholkopf, B. 2001. An introduction to kernel-based learning algorithms. *IEEE transactions on neural networks* 12(2):181–201.
- [26] North, M.; Collier, N.; and Vos, J. 2006. Experiences creating three implementations of the repast agent modeling toolkit. *ACM Transactions on Modeling and Computer Simulation (TOMACS)* 16(1):25.
- [27] Parrish, J.; Viscido, S.; and Grunbaum, D. 2002. Self-organized fish schools: an examination of emergent properties. *Biological Bulletin, Marine Biological Laboratory, Woods Hole* 202(3):296–305.

- [28] Parunak, H.; Savit, R.; and Riolo, R. 1998. Agent-Based Modeling vs Equation-Based Modeling: A Case Study and Users' Guide. *Lecture Notes in Computer Science* 1534:10–25.
- [29] Patrikalakis, N.; Maekawa, T.; Ko, K.; and Mukundan, H. 1993. Surface-to-surface intersections. *IEEE Computer Graphics and Applications* 13(1):89–95.
- [30] Pratt, S.; Sumpter, D.; Mallon, E.; and Franks, N. 2005. An agent-based model of collective nest choice by the ant *Temnothorax albipennis*. *Animal Behaviour* 70(5):1023–1036.
- [31] Reynolds, C. W. 1987. Flocks, herds, and schools: A distributed behavioral model. In *Computer Graphics, 21(4) (SIGGRAPH '87 Conference Proceedings)*, 25–34.
- [32] Reynolds, C. 1999. Steering behaviors for autonomous characters. In *Game Developers Conference*, volume 1999.
- [33] Schelhorn, T.; O'Sullivan, D.; Haklay, M.; and Thurstain-Goodwin, M. 1999. STREETS: an agent-based pedestrian model. In *Proceedings of Computers in Urban Planning and Urban Management*.
- [34] Shi, Y., and Eberhart, R. 1998. Parameter selection in particle swarm optimization. *Lecture notes in computer science* 591–600.
- [35] Smola, A., and Schölkopf, B. 2004. A tutorial on support vector regression. *Statistics and Computing* 14(3):199–222.
- [36] Spears, W.; Spears, D.; Hamann, J.; and Heil, R. 2004. Distributed, physics-based control of swarms of vehicles. *Autonomous Robots* 17(2):137–162.
- [37] Tisue, S., and Wilensky, U. 2004. NetLogo: A simple environment for modeling complexity. In *International Conference on Complex Systems*, 16–21.

- [38] Van den Bergh, F., and Engelbrecht, A. 2006. A study of particle swarm optimization particle trajectories. *Information Sciences* 176(8):937–971.
- [39] Wilensky, U. 1997a. NetLogo Flocking model. Technical report, Northwestern University.
- [40] Wilensky, U. 1997b. NetLogo Traffic Basic model. Technical report, Northwestern University.
- [41] Wilensky, U. 1997c. NetLogo Wolf Sheep Predation model. Technical report, Northwestern University.
- [42] Zambaka, C.; HandUber, J.; and Saunders-Newton, D. 2009. Modeling and Simulating Community Sentiments and Interactions at the Pacific Missile Range Facility.